

# gRPC Go

## GothamGo 2015

Sameer Ajmani  
Tech Lead Manager, Go team, Google

# Video

This talk was presented at GothamGo in New York City, October 2015.

[Watch the talk on YouTube](https://www.youtube.com/watch?v=vTlyz2QfExc&index=7&list=PLeGxIOPLk9ELh9tsPZMzau6CzMjfiMzp9-) (https://www.youtube.com/watch?v=vTlyz2QfExc&index=7&list=PLeGxIOPLk9ELh9tsPZMzau6CzMjfiMzp9-)

# RPC isn't just Remote Procedure Call

In Go, an RPC **starts a goroutine** running on the server and provides *message passing* between the client and server goroutines.

**Unary RPC:** the client sends a *request* to the server, then the server sends a *response*.

**Streaming RPC:** the client and server may each send one or more messages.

An RPC ends when:

- both sides are done sending messages
- either side disconnects
- the RPC is canceled or times out

This talk will show how we connect RPCs and streams with goroutines and channels.

## Unary RPC: one request, one response

Example: a mobile Maps app requests a route from point A to point B.

On the client side, an RPC blocks until it's done or canceled.

A client uses multiple goroutines to run many RPCs simultaneously.

Each RPC is an exchange between a client goroutine and a server goroutine.

# Streaming RPC provides bidirectional message-passing

A client starts a stream with a server.

Messages sent on a stream are delivered FIFO.

Many streams can run simultaneously between the same client and server.

The transport provides buffering and flow control.

Examples:

- bidirectional stream: chat session
- server → client stream: stock ticker
- client → server stream: sensor aggregation

# gRPC is a new RPC system from Google

[grpc.io](http://grpc.io) (<http://grpc.io>)

Provides RPC and streaming RPC

Ten languages: **C, Java, Go**, C++, Node.js, Python, Ruby, Objective-C, PHP, and C#

IDL: **Proto3**

Transport: **HTTP2**

[golang.org/x/net/context](http://golang.org/x/net/context) (<http://golang.org/x/net/context>) for deadlines, cancelation, and request-scoped values

[golang.org/x/net/trace](http://golang.org/x/net/trace) (<http://golang.org/x/net/trace>) for real-time request traces and connection logging

# gRPC users

150+ imports of [google.golang.org/grpc](https://godoc.org/google.golang.org/grpc) (<https://godoc.org/google.golang.org/grpc?importers>) on [godoc.org](http://godoc.org) (<http://godoc.org>)

- [Apcera/Kurma](https://github.com/apcera/kurma) (<https://github.com/apcera/kurma>): container OS
- [Bazil](http://bazil.org) (<http://bazil.org>): distributed file system
- [CoreOS/Etcd](http://coreos.com/etcd/) (<http://coreos.com/etcd/>): distributed consistent key-value store
- [Google Cloud Bigtable](https://godoc.org/google.golang.org/cloud/bigtable) (<https://godoc.org/google.golang.org/cloud/bigtable>): sparse table storage
- [Monetas/Bitmessage](https://github.com/monetas/bmd) (<https://github.com/monetas/bmd>): transaction platform
- [Pachyderm](http://www.pachyderm.io/) (<http://www.pachyderm.io/>): containerized data analytics
- [YouTube/Vitess](http://vitess.io/) (<http://vitess.io/>): storage platform for scaling MySQL

# Demos and Code: Google search



# Protocol definition

```
syntax = "proto3";

service Google {
    // Search returns a Google search result for the query.
    rpc Search(Request) returns (Result) {
    }
}

message Request {
    string query = 1;
}

message Result {
    string title = 1;
    string url = 2;
    string snippet = 3;
}
```

# Generated code

```
protoc ./search.proto --go_out=plugins=grpc:.
```


```
type GoogleClient interface {  
    // Search returns a Google search result for the query.  
    Search(ctx context.Context, in *Request, opts ...grpc.CallOption) (*Result, error)  
}
```

```
type GoogleServer interface {  
    // Search returns a Google search result for the query.  
    Search(context.Context, *Request) (*Result, error)  
}
```


```
type Request struct {  
    Query string `protobuf:"bytes,1,opt,name=query" json:"query,omitempty"`  
}
```

```
type Result struct {  
    Title   string `protobuf:"bytes,1,opt,name=title" json:"title,omitempty"`  
    Url     string `protobuf:"bytes,2,opt,name=url" json:"url,omitempty"`  
    Snippet string `protobuf:"bytes,3,opt,name=snippet" json:"snippet,omitempty"`  
}
```

# System diagram




# Frontend runs Search on both backends and returns first result



## Demo client --mode=search

- Frontend request traces
- Backend request traces
- Connection event logs

# Client code



# Client code (main)

```
import pb "golang.org/x/talks/2015/gotham-grpc/search"
```

```
func main() {  
    flag.Parse()  
  
    // Connect to the server.  
    conn, err := grpc.Dial(*server, grpc.WithInsecure())  
    if err != nil {  
        log.Fatalf("fail to dial: %v", err)  
    }  
    defer conn.Close()  
    client := pb.NewGoogleClient(conn)  
  
    // Run the RPC.  
    switch *mode {  
    case "search":  
        search(client, *query)  
    case "watch":  
        watch(client, *query)  
    default:  
        log.Fatalf("unknown mode: %q", *mode)  
    }  
}
```

## Client code (search)


```
func search(client pb.GoogleClient, query string) {  
    ctx, cancel := context.WithTimeout(context.Background(), 80*time.Millisecond)  
    defer cancel()  
    req := &pb.Request{Query: query}  
    res, err := client.Search(ctx, req)  
    if err != nil {  
        log.Fatal(err)  
    }  
    fmt.Println(res)  
}
```

RPCs block but can be canceled using a Context.

gRPC propagates cancelation from client to server.



# Backend code



## Backend code (main)

```
lis, err := net.Listen("tcp", fmt.Sprintf(":%d", 36061+*index)) // RPC port
if err != nil {
    log.Fatalf("failed to listen: %v", err)
}
g := grpc.NewServer()
pb.RegisterGoogleServer(g, new(server))
g.Serve(lis)
```

`new(server)` must implement the `GoogleServer` interface:

```
type GoogleServer interface {
    // Search returns a Google search result for the query.
    Search(context.Context, *Request) (*Result, error)
    // Watch returns a stream of Google search results for the query.
    Watch(*Request, Google_WatchServer) error
}
```

Each call to `Search` or `Watch` runs in its own goroutine.

## Backend code (search)


ctx.Done is closed when the RPC is canceled, times out, or returns:

```
func (s *server) Search(ctx context.Context, req *pb.Request) (*pb.Result, error) {
    d := randomDuration(100 * time.Millisecond)
    logSleep(ctx, d)
    select {
    case <-time.After(d):
        return &pb.Result{
            Title: fmt.Sprintf("result for [%s] from backend %d", req.Query, *index),
        }, nil
    case <-ctx.Done():
        return nil, ctx.Err()
    }
}
```

If tracing is enabled, log the sleep duration:

```
func logSleep(ctx context.Context, d time.Duration) {
    if tr, ok := trace.FromContext(ctx); ok {
        tr.LazyPrintf("sleeping for %s", d)
    }
}
```

# Frontend code



## Frontend code (search)

Search returns as soon as it gets the first result.

gRPC cancels the remaining backend. Search RPCs by via ctx:

```
func (s *server) Search(ctx context.Context, req *pb.Request) (*pb.Result, error) {
    c := make(chan result, len(s.backends))
    for _, b := range s.backends {
        go func(backend pb.GoogleClient) {
            res, err := backend.Search(ctx, req)
            c <- result{res, err}
        }(b)
    }
    first := <-c
    return first.res, first.err
}
```

```
type result struct {
    res *pb.Result
    err error
}
```

# Streaming RPC

# Add Watch to the Google service

```
syntax = "proto3";

service Google {
    // Search returns a Google search result for the query.
    rpc Search(Request) returns (Result) {
    }
    // Watch returns a stream of Google search results for the query.
    rpc Watch(Request) returns (stream Result) {
    }
}

message Request {
    string query = 1;
}

message Result {
    string title = 1;
    string url = 2;
    string snippet = 3;
}
```

# Generated code

```
type GoogleClient interface {  
    // Search returns a Google search result for the query.  
    Search(ctx context.Context, in *Request, opts ...grpc.CallOption) (*Result, error)  
    // Watch returns a stream of Google search results for the query.  
    Watch(ctx context.Context, in *Request, opts ...grpc.CallOption) (Google_WatchClient, error)  
}
```

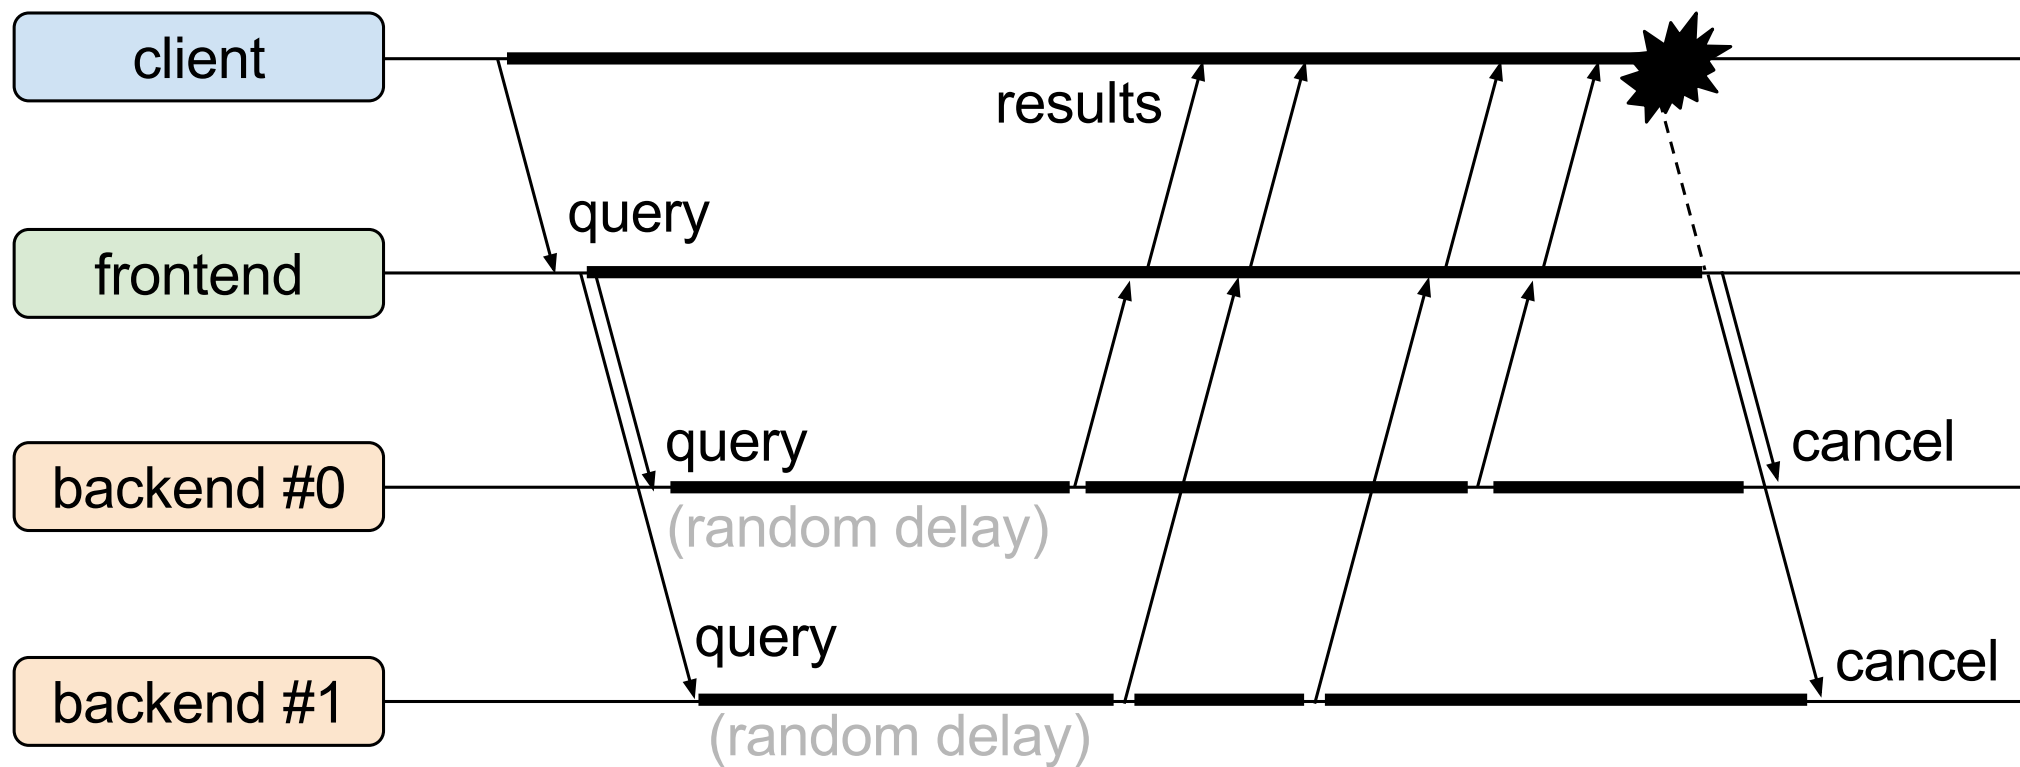
```
type GoogleServer interface {  
    // Search returns a Google search result for the query.  
    Search(context.Context, *Request) (*Result, error)  
    // Watch returns a stream of Google search results for the query.  
    Watch(*Request, Google_WatchServer) error  
}
```

```
type Google_WatchClient interface {  
    Recv() (*Result, error)  
    grpc.ClientStream  
}
```

```
type Google_WatchServer interface {  
    Send(*Result) error  
    grpc.ServerStream  
}
```




# Frontend runs Watch on both backends and merges results



## Demo client --mode=watch

- Active stream traces
- Cancelation


# Client code



# Client code (watch)

```
func watch(client pb.GoogleClient, query string) {  
    ctx, cancel := context.WithCancel(context.Background())  
    defer cancel()  
    req := &pb.Request{Query: query}  
    stream, err := client.Watch(ctx, req)  
    if err != nil {  
        log.Fatal(err)  
    }  
    for {  
        res, err := stream.Recv()  
        if err == io.EOF {  
            fmt.Println("and now your watch is ended")  
            return  
        }  
        if err != nil {  
            log.Fatal(err)  
        }  
        fmt.Println(res)  
    }  
}
```


# Backend code



## Backend code (watch)

```
func (s *server) Watch(req *pb.Request, stream pb.Google_WatchServer) error {
    ctx := stream.Context()
    for i := 0; ; i++ {
        d := randomDuration(1 * time.Second)
        logSleep(ctx, d)
        select {
        case <-time.After(d):
            err := stream.Send(&pb.Result{
                Title: fmt.Sprintf("result %d for [%s] from backend %d", i, req.Query, *index),
            })
            if err != nil {
                return err
            }
        case <-ctx.Done():
            return ctx.Err()
        }
    }
}
```

# Frontend code



## Frontend code (watch)

```
func (s *server) Watch(req *pb.Request, stream pb.Google_WatchServer) error {
    ctx := stream.Context()
    c := make(chan result)
    var wg sync.WaitGroup
    for _, b := range s.backends {
        wg.Add(1)
        go func(backend pb.GoogleClient) {
            defer wg.Done()
            watchBackend(ctx, backend, req, c)
        }(b)
    }
    go func() {
        wg.Wait()
        close(c)
    }()
    for res := range c {
        if res.err != nil {
            return res.err
        }
        if err := stream.Send(res.res); err != nil {
            return err
        }
    }
    return nil
}
```



## Frontend code (watchBackend)

Watch returns on first error; this closes `ctx.Done` and signals `watchBackend` to exit.

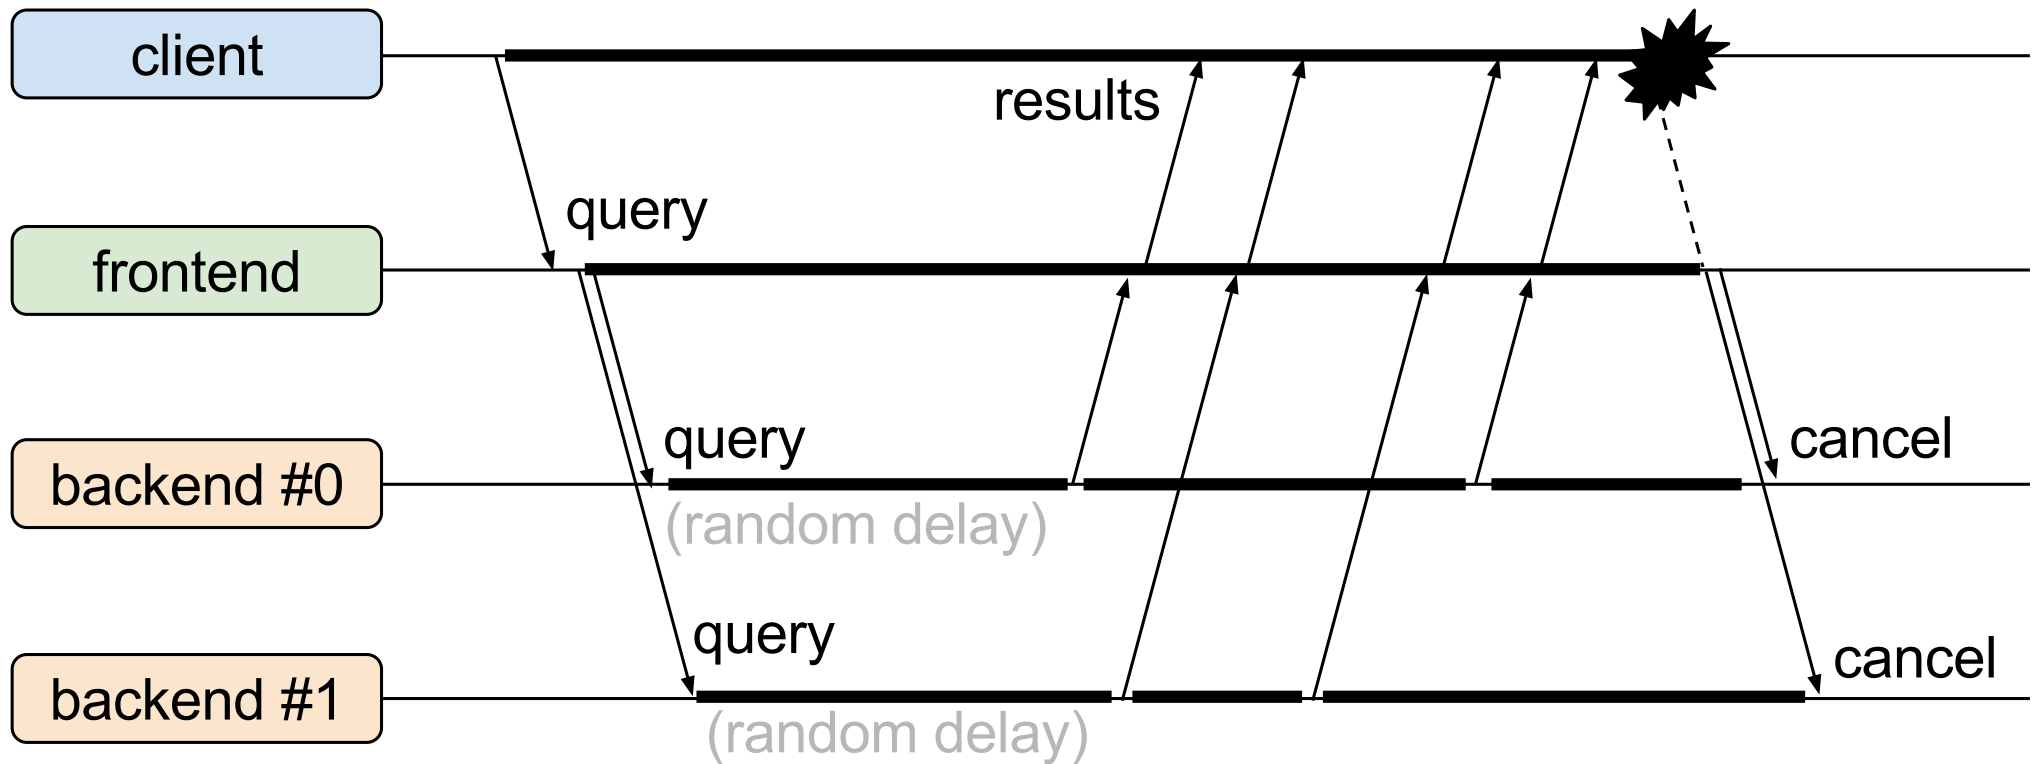
```
func watchBackend(ctx context.Context, backend pb.GoogleClient, req *pb.Request, c chan<- result) {  
    stream, err := backend.Watch(ctx, req)  
    if err != nil {  
        select {  
        case c <- result{err: err}:  
        case <-ctx.Done():  
        }  
        return  
    }  
    for {  
        res, err := stream.Recv()  
        select {  
        case c <- result{res, err}:  
            if err != nil {  
                return  
            }  
        case <-ctx.Done():  
            return  
        }  
    }  
}
```



# gRPC extends the Go programming model over the network

Go gRPC works smoothly with goroutines, channels, and cancelation.

It is an excellent fit for building parallel, distributed, and streaming systems.



# References

- [grpc.io](http://grpc.io) (<http://grpc.io>) - gRPC reference and tutorials
- [github.com/golang/protobuf](https://github.com/golang/protobuf) (<https://github.com/golang/protobuf>) - Protocol buffers
- [golang.org/x/net/http2](http://golang.org/x/net/http2) (<http://golang.org/x/net/http2>) - HTTP2
- [golang.org/x/net/trace](http://golang.org/x/net/trace) (<http://golang.org/x/net/trace>) - Request traces and event logs
- [golang.org/x/net/context](http://golang.org/x/net/context) (<http://golang.org/x/net/context>) - Cancellation and request-scoped data
- [blog.golang.org/pipelines](http://blog.golang.org/pipelines) (<http://blog.golang.org/pipelines>) - Streaming data pipelines

**Thanks to** Qi Zhao, David Symonds, Brad Fitzpatrick, and the rest.

## Questions?

Sameer Ajmani

Tech Lead Manager, Go team, Google

[@Sajma](https://twitter.com/Sajma) ([twitter.com/Sajma](https://twitter.com/Sajma))

[sameer@golang.org](mailto:sameer@golang.org) (<mailto:sameer@golang.org>)

# Thank you

Sameer Ajmani

Tech Lead Manager, Go team, Google

[@Sajma](http://twitter.com/Sajma) (<http://twitter.com/Sajma>)

[sameer@golang.org](mailto:sameer@golang.org) (<mailto:sameer@golang.org>)

